

به نام پیگانه معبود بخشنده مهربان

**مبانی یادگیری ماشین**

# **Machine Learning Foundations**

**گروه هوش مصنوعی، دانشکده مهندسی کامپیوتر، دانشگاه اصفهان**

**ترم اول سال تحصیلی ۰۳-۰۲**

**ارائه دهنده : پیمان ادیبی**

---

دسته‌بند رگرسیون منطقی

Logistic Regression Classifier

---


# رگرسیون منطقی (یک دسته بند دودویی)

- An alternative: replace the  $\text{sign}(\cdot)$  with the **sigmoid** or **logistic function**
- We assumed a particular functional form: sigmoid applied to a linear function of the data

$$y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

where the sigmoid is defined as

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



- The output is a smooth function of the inputs and the weights. It can be seen as a smoothed and differentiable alternative to  $\text{sign}(\cdot)$

# رگرسیون منطقی (یک دسته بند دودویی)

- We assumed a particular functional form: sigmoid applied to a linear function of the data

$$y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

where the sigmoid is defined as

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$


- ▶ One parameter per data dimension (feature) and the bias
- ▶ Features can be discrete or continuous
- ▶ Output of the model: value  $y \in [0, 1]$
- ▶ Allows for gradient-based learning of the parameters

# شكل تابع سيغمويدي


- Let's look at how modifying  $w$  changes the shape of the function
- 1D example:

$$y = \sigma(w_1 x + w_0)$$


$$w_0 = 0, w_1 = 1$$



$$w_0 = 0, w_1 = 0.5$$



$$w_0 = -2, w_1 = 1$$



- Demo



✓380

# رگرسیون منطقی - تفسیر احتمالاتی

- If we have a value between 0 and 1, let's use it to model class probability

$$p(C = 0|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0) \quad \text{with} \quad \sigma(z) = \frac{1}{1 + \exp(-z)}$$

- Substituting we have

$$p(C = 0|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)}$$

- Suppose we have two classes, how can I compute  $p(C = 1|\mathbf{x})$ ?
- Use the marginalization property of probability



$$p(C = 1|\mathbf{x}) + p(C = 0|\mathbf{x}) = 1$$

- Thus (show matlab)

$$p(C = 1|\mathbf{x}) = 1 - \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)} = \frac{\exp(-\mathbf{w}^T \mathbf{x} - w_0)}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)}$$

# رگرسیون منطقی - مرز تصمیم

- What is the decision boundary for logistic regression?
- $p(C = 1|\mathbf{x}, \mathbf{w}) = p(C = 0|\mathbf{x}, \mathbf{w}) = 0.5$
- $p(C = 0|\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0) = 0.5$ , where  $\sigma(z) = \frac{1}{1+\exp(-z)}$
- Decision boundary:  $\mathbf{w}^T \mathbf{x} + w_0 = 0$
- Logistic regression has a **linear** decision boundary






# مثال

- **Problem:** Given the number of hours a student spent learning, will (s)he pass the exam?
- Training data (top row:  $x^{(i)}$ , bottom row:  $t^{(i)}$ )

Hours	0.50	0.75	1.00	1.25	1.50	1.75	1.75	2.00	2.25	2.50	2.75	3.00	3.25	3.50	4.00	4.25	4.50	4.75	5.00	5.50
Pass	0	0	0	0	0	0	1	0	1	0	1	0	1	0	1	1	1	1	1	1

- Learn  $w$  for our model, i.e. logistic regression (coming up)
- Make predictions:



Hours of study	Probability of passing exam
1	0.07
2	0.26
3	0.61
4	0.87
5	0.97

# رگرسیون منطقی - یادگیری

- How should we learn the weights  $\mathbf{w}$ ,  $w_0$ ?
- We have a probabilistic model
- Let's use maximum likelihood

(simplify notation: we will write  $\mathbf{w}$  to represent both  $\mathbf{w}$  and  $w_0$ )

# تابع شباهت شرطی

- Assume  $t \in \{0, 1\}$ , we can write the probability distribution of each of our training points  $p(t^{(1)}, \dots, t^{(N)} | \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}; \mathbf{w})$
- Assuming that the training examples are **sampled IID**: independent and identically distributed, we can write the *likelihood function*:

$$L(\mathbf{w}) = p(t^{(1)}, \dots, t^{(N)} | \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}; \mathbf{w}) = \prod_{i=1}^N p(t^{(i)} | \mathbf{x}^{(i)}; \mathbf{w})$$

- We can write each probability as (will be useful later):

$$\begin{aligned} p(t^{(i)} | \mathbf{x}^{(i)}; \mathbf{w}) &= p(C = 1 | \mathbf{x}^{(i)}; \mathbf{w})^{t^{(i)}} p(C = 0 | \mathbf{x}^{(i)}; \mathbf{w})^{1-t^{(i)}} \\ &= \left(1 - p(C = 0 | \mathbf{x}^{(i)}; \mathbf{w})\right)^{t^{(i)}} p(C = 0 | \mathbf{x}^{(i)}; \mathbf{w})^{1-t^{(i)}} \end{aligned}$$

- We can learn the model by **maximizing the likelihood**

$$\max_{\mathbf{w}} L(\mathbf{w}) = \max_{\mathbf{w}} \prod_{i=1}^N p(t^{(i)} | \mathbf{x}^{(i)}; \mathbf{w})$$

- Easier to maximize the log likelihood  $\log L(\mathbf{w})$

# تابع اتلاف

$$\begin{aligned} L(\mathbf{w}) &= \prod_{i=1}^N p(t^{(i)}|\mathbf{x}^{(i)}) \quad (\text{likelihood}) \\ &= \prod_{i=1}^N \left(1 - p(C=0|\mathbf{x}^{(i)})\right)^{t^{(i)}} p(C=0|\mathbf{x}^{(i)})^{1-t^{(i)}} \end{aligned}$$

- We can convert the maximization problem into minimization so that we can write the **loss function**:

$$\begin{aligned} \ell_{\log}(\mathbf{w}) &= -\log L(\mathbf{w}) \\ &= -\sum_{i=1}^N \log p(t^{(i)}|\mathbf{x}^{(i)}; \mathbf{w}) \\ &= -\sum_{i=1}^N t^{(i)} \log(1 - p(C=0|\mathbf{x}^{(i)}, \mathbf{w})) - \sum_{i=1}^N (1 - t^{(i)}) \log p(C=0|\mathbf{x}^{(i)}; \mathbf{w}) \end{aligned}$$

- Is there a closed form solution?
- It's a convex function of  $\mathbf{w}$ . Can we get the global optimum?

# نزول در راستای گرادیان

$$\min_{\mathbf{w}} \ell(\mathbf{w}) = \min_{\mathbf{w}} \left\{ - \sum_{i=1}^N t^{(i)} \log(1 - p(C = 0 | \mathbf{x}^{(i)}, \mathbf{w})) - \sum_{i=1}^N (1 - t^{(i)}) \log p(C = 0 | \mathbf{x}^{(i)}, \mathbf{w}) \right\}$$

- Gradient descent: iterate and at each iteration compute steepest direction towards optimum, move in that direction, step-size  $\lambda$

$$w_j^{(t+1)} \leftarrow w_j^{(t)} - \lambda \frac{\partial \ell(\mathbf{w})}{\partial w_j}$$

- You can write this in vector form

$$\nabla \ell(\mathbf{w}) = \left[ \frac{\partial \ell(\mathbf{w})}{\partial w_0}, \dots, \frac{\partial \ell(\mathbf{w})}{\partial w_k} \right]^T, \quad \text{and} \quad \Delta(\mathbf{w}) = -\lambda \nabla \ell(\mathbf{w})$$

- But where is  $\mathbf{w}$ ?

$$p(C = 0 | \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)}, \quad p(C = 1 | \mathbf{x}) = \frac{\exp(-\mathbf{w}^T \mathbf{x} - w_0)}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)}$$



# نزول در راستای گرادیان - محاسبات

- The loss is

$$\ell_{\log\text{-loss}}(\mathbf{w}) = -\sum_{i=1}^N t^{(i)} \log p(C = 1|\mathbf{x}^{(i)}, \mathbf{w}) - \sum_{i=1}^N (1-t^{(i)}) \log p(C = 0|\mathbf{x}^{(i)}, \mathbf{w})$$

where the probabilities are

$$p(C = 0|\mathbf{x}, \mathbf{w}) = \frac{1}{1 + \exp(-z)} \quad p(C = 1|\mathbf{x}, \mathbf{w}) = \frac{\exp(-z)}{1 + \exp(-z)}$$

and  $z = \mathbf{w}^T \mathbf{x} + w_0$

- We can simplify

$$\begin{aligned} \ell(\mathbf{w})_{\log\text{-loss}} &= \sum_i t^{(i)} \log(1 + \exp(-z^{(i)})) + \sum_i t^{(i)} z^{(i)} + \sum_i (1 - t^{(i)}) \log(1 + \exp(-z^{(i)})) \\ &= \sum_i \log(1 + \exp(-z^{(i)})) + \sum_i t^{(i)} z^{(i)} \end{aligned}$$

- Now it's easy to take derivatives

# نزول در راستای گرادیان - محاسبات

$$\ell(\mathbf{w}) = \sum_i t^{(i)} z^{(i)} + \sum_i \log(1 + \exp(-z^{(i)}))$$

- Now it's easy to take derivatives
- Remember  $z = \mathbf{w}^T \mathbf{x} + w_0$

$$\frac{\partial \ell}{\partial w_j} = \sum_i \left( t^{(i)} x_j^{(i)} - x_j^{(i)} \cdot \frac{\exp(-z^{(i)})}{1 + \exp(-z^{(i)})} \right)$$

- What's  $x_j^{(i)}$ ? The  $j$ -th dimension of the  $i$ -th training example  $\mathbf{x}^{(i)}$
- And simplifying

$$\frac{\partial \ell}{\partial w_j} = \sum_i x_j^{(i)} \left( t^{(i)} - p(C = 1 | \mathbf{x}^{(i)}; \mathbf{w}) \right)$$

- Don't get confused with indices:  $j$  for the weight that we are updating and  $i$  for the training example

# نزول در راستای گرادیان - بهنگام سازی

- Putting it all together (plugging the update into gradient descent):

Gradient descent for logistic regression:

$$w_j^{(t+1)} \leftarrow w_j^{(t)} - \lambda \sum_i x_j^{(i)} \left( t^{(i)} - p(C = 1 | \mathbf{x}^{(i)}; \mathbf{w}) \right)$$

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \lambda \sum_i \mathbf{x}^{(i)} \left( t^{(i)} - p(C = 1 | \mathbf{x}^{(i)}; \mathbf{w}) \right)$$



where:

$$p(C = 1 | \mathbf{x}^{(i)}; \mathbf{w}) = \frac{\exp(-\mathbf{w}^T \mathbf{x} - w_0)}{1 + \exp(-\mathbf{w}^T \mathbf{x} - w_0)} = 1 - \frac{1}{1 + \exp(\mathbf{w}^T \mathbf{x} + w_0)}$$

- This is all there is to learning in logistic regression. Simple, huh?




# منحنی یادگیری - انتخاب نرخ یادگیری





# منحنی یادگیری - انتخاب نرخ یادگیری

What about a larger step-size?

Can cause divergence!



# رگرسیون منطقی VS رگرسیون کمترین مربعات



✓414

# رگرسیون منطقی - تنظیم

- We can also look at  $p(\mathbf{w}|\{t\}, \{\mathbf{x}\}) \propto p(\{t\}|\{\mathbf{x}\}, \mathbf{w}) p(\mathbf{w})$   
with  $\{t\} = (t^{(1)}, \dots, t^{(N)})$ , and  $\{\mathbf{x}\} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)})$

$$\max_{\mathbf{w}} \log \left[ \underbrace{p(\mathbf{w})}_{\text{prior on parameters } \mathbf{w}} \prod_i p(t^{(i)}|\mathbf{x}^{(i)}, \mathbf{w}) \right]$$

- This is a form of regularization
- Helps avoid large weights and **overfitting**
- For example, define prior: normal distribution, zero mean and identity covariance  $p(\mathbf{w}) = \mathcal{N}(0, \alpha^{-1}\mathbf{I})$
- This prior pushes parameters towards zero
- Including this prior the new gradient is

$$w_j^{(t+1)} \leftarrow w_j^{(t)} - \lambda \frac{\partial \ell(\mathbf{w})}{\partial w_j} - \lambda \alpha w_j^{(t)}$$

- How do we decide the best value of  $\alpha$  (or a hyper-parameter in general)?

# اعتبار سنجی (متقابل)

Tuning hyper-parameters:

- **Never use test data for tuning the hyper-parameters**
- We can divide the set of training examples into two disjoint sets: **training** and **validation**
- Use the first set (i.e., training) to estimate the weights  $\mathbf{w}$  for different values of  $\alpha$
- Use the second set (i.e., validation) to estimate the best  $\alpha$ , by evaluating how well the classifier does on this second set
- This tests how well it generalizes to unseen data

Leave-p-out cross-validation

Leave-1-out cross-validation

k-fold cross-validation





### Advantages:

- Easily extended to multiple classes (thoughts?)
- Natural probabilistic view of class predictions
- Quick to train
- Fast at classification
- Good accuracy for many simple data sets
- Resistant to overfitting
- Can interpret model coefficients as indicators of feature importance

### Less good:

- Linear decision boundary (too simple for more complex problems?)

[Slide by: Jeff Howbert]